


Vasil Sarafov

David Markvica

Stefan Brunthaler


Fuzzers have limitations

- Bug finding is undecidable, fuzzers rely on heuristics.
- How hard is “too hard” for a fuzzer?
- Which heuristics work best for specific obstacles?






How to identify fuzzer limitations?

- TEPHRA: principled methodology that yields optimistic upper bounds.
- Reduces overfitting & complements existing evaluation approaches.
- Systematically probes the semantic space.



31 fuzzers, 26 C/C++ semantic obstacles, 37.4 CPU years



Fuzzer	Chain										Interval 10 ⁻⁶						
	bool	Unsigned Int			Signed Int			Float		Cmp	Unsigned Int		Signed Int		Float		
		8	16	32	64	8	16	32	64	str	mem	16	32	64	32	64	
AFL GCC	15	3	1	0	0	3	1	0	0	0	0	2	0	3	14	0	
AFL clang-fast	14	15	1	0	0	14	1	0	0	0	0	2	0	3	14	0	
AFL libtokencap	16	15	1	0	0	15	1	0	0	0	0	2	0	3	13	0	
AFL QEMU	32	4	1	0	0	4	1	0	0	0	0	0	2	0	3	14	0
AFL++	15	15	2	64	32	14	5	3	32	0	0	5	32	0	3	13	0
AFL++ Mopt	15	14	0	32	0	14	5	3	32	0	0	0	32	0	3	13	0
AFL++ CmpLog	15	3	2	0	0	3	2	0	0	0	0	0	2	0	3	13	0
AFL++ QEMU CmpCov	16	15	7	3	5	15	7	3	5	0	0	0	8	0	0	0	0
AFLGo	15	3	2	0	0	2	1	0	0	0	0	0	2	0	3	13	0
Angora	4096	8192	4096	2048	1024	8192	4096	2048	1024	0	0	0	2	0	1	11	0
Angora mb	8192	8192	4096	2048	1024	8192	4096	2048	1024	0	0	0	2	0	11	2	-1
Angora random	8192	8192	10	0	0	64	3	14	0	0	0	0	2	0	3	13	-3
DARWIN	512	64	1	0	0	64	3	0	0	0	0	0	2	0	3	13	0
dataAFLow	14	2	1	0	0	2	1	0	0	0	0	0	2	0	3	13	0
DDFuzz	512	16	2	4	10	32	1	10	10	0	0	0	16	32	0	3	13
dev-urandom	13	3	1	0	0	3	1	0	0	0	0	0	3	12	0	2	12
EcoFuzz	128	14	1	0	0	15	7	9	0	0	0	0	4	0	7	0	0
FA-Fuzz	512	14	2	0	0	13	1	0	0	0	0	0	2	0	3	13	0
FairFuzz	1024	128	7	0	0	128	8	0	0	0	0	0	2	0	3	13	0
honggfuzz	256	32	32	16	16	32	32	32	16	0	0	0	32	32	0	0	-1
honggfuzz QEMU	16	13	6	3	1	7	3	2	0	0	0	0	1	13	0	8	-4
KLEE	128	128	64	32	16	128	64	32	16	0	0	0	16	128	0	0	0
LibAFL	16	7	3	16	15	7	10	16	0	0	0	5	32	0	3	12	-5
LibAFL-libFuzzer	512	64	5	1024	64	3	1024	1024	0	0	0	3	0	0	0	0	-5
libfuzzer	512	256	64	128	128	32	64	64	0	0	0	64	64	0	0	0	-3
libfuzzer Entropic	512	256	32	16	64	256	32	128	64	0	0	64	64	0	0	0	-5
Radamsa	13	2	1	0	2	1	0	0	0	0	0	0	2	4	13	-4	-3
SymCC	13	9	9	9	10	9	9	10	0	0	0	0	16384	0	0	0	0
Wingfuzz	1024	256	128	32	128	256	32	256	128	0	0	64	64	0	0	0	-3

Results

- All fuzzers struggle with certain semantic constructs.
- Support for rational numbers and character strings lacking.
- Signed integers more difficult than unsigned.
- Overtuning for 32- and 64-bit types, neglecting 8- and 16-bit.
- No fuzzer excels across all obstacles.
- A single obstacle can degrade overall performance.
- ~90% reliable bypasses; ~10% due to fluctuating randomness.

Next Steps

- Further experiments.
- Fuzzing based on an obstacle profile.
- Extend implementation with additional semantics and PLs.

